

# Predicting determinants of crimes in US Communities

This project uses the communities dataset from UCI learning repository to explain factors responsible for crimes in many US communities. The data contains 2215 instances with 147 records. We have carefully selected about 20 variables which we feel can trigger crimes in the communities.

Some variables are constructed in different ways to essentially measure the same thing. For instance, reporting number of people under poverty and percentage of people living under poverty, the 2 variables essentially are quantifying the same thing, in such situations, we stick to only 1 of the variables, preferably percentage measures.

## Importing the necessary libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

## Import the dataset

```
In [2]: data = pd.read_csv("/content/drive/MyDrive/Crime Analytics/crimenew.csv")
data.head()
```

Out[2]:

|   | householdsize | racepctblack | racePctWhite | racePctAsian | racePctHisp | agePct12t29 | pctWWage | pctWSocSec | pctWPubAsst | perCapInc | ... | MalePctNevMz |
|---|---------------|--------------|--------------|--------------|-------------|-------------|----------|------------|-------------|-----------|-----|--------------|
| 0 | 2.57          | 0.18         | 98.48        | 0.84         | 1.66        | 24.90       | 77.02    | 35.90      | 5.43        | 16201     | ... | 31.          |
| 1 | 2.40          | 6.50         | 91.93        | 1.13         | 1.17        | 27.58       | 79.77    | 27.82      | 6.00        | 15138     | ... | 33.          |
| 2 | 2.92          | 1.41         | 96.01        | 1.76         | 3.61        | 27.30       | 89.65    | 15.44      | 3.02        | 17642     | ... | 28.          |
| 3 | 2.59          | 1.58         | 90.57        | 7.00         | 5.08        | 23.23       | 81.68    | 31.65      | 2.30        | 19643     | ... | 30.          |
| 4 | 2.38          | 0.42         | 99.19        | 0.24         | 0.52        | 23.42       | 66.11    | 44.02      | 6.67        | 10246     | ... | 26.          |

5 rows × 29 columns



```
In [3]: data.tail()
```

Out[3]:

|      | householdsize | racepctblack | racePctWhite | racePctAsian | racePctHisp | agePct12t29 | pctWWage | pctWSocSec | pctWPubAsst | perCapInc | ... | MalePctNe |
|------|---------------|--------------|--------------|--------------|-------------|-------------|----------|------------|-------------|-----------|-----|-----------|
| 2210 | 2.76          | 0.82         | 97.60        | 1.26         | 1.86        | 24.96       | 78.41    | 30.55      | 2.55        | 26895     | ... |           |
| 2211 | 2.50          | 7.58         | 81.61        | 1.79         | 16.27       | 30.11       | 74.42    | 31.61      | 9.08        | 14715     | ... |           |
| 2212 | 2.72          | 6.57         | 86.75        | 3.94         | 7.69        | 28.95       | 83.63    | 23.55      | 3.98        | 19300     | ... |           |
| 2213 | 2.63          | 2.13         | 93.26        | 3.49         | 4.42        | 22.08       | 80.52    | 28.24      | 2.38        | 46070     | ... |           |
| 2214 | 2.65          | 1.52         | 96.79        | 1.03         | 2.31        | 24.75       | 82.47    | 28.50      | 3.14        | 19099     | ... |           |

5 rows × 29 columns



```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2215 entries, 0 to 2214
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype
---  -
0   householdsize          2215 non-null   float64
1   racepctblack            2215 non-null   float64
2   racePctWhite            2215 non-null   float64
3   racePctAsian            2215 non-null   float64
4   racePctHisp             2215 non-null   float64
5   agePct12t29             2215 non-null   float64
6   pctWWage                2215 non-null   float64
7   pctWSocSec              2215 non-null   float64
8   pctWPubAsst             2215 non-null   float64
9   perCapInc               2215 non-null   int64
10  whitePerCap             2215 non-null   int64
11  blackPerCap             2215 non-null   int64
```

```

12 indianPerCap      2215 non-null int64
13 AsianPerCap       2215 non-null int64
14 OtherPerCap       2214 non-null float64
15 HispPerCap        2215 non-null int64
16 PctPopUnderPov    2215 non-null float64
17 PctNotHSGrad      2215 non-null float64
18 PctUnemployed     2215 non-null float64
19 MalePctNevMarr    2215 non-null float64
20 TotalPctDiv       2215 non-null float64
21 PctKidsBornNeverMar 2215 non-null float64
22 PctImmigRec5      2215 non-null float64
23 PctPersDenseHous  2215 non-null float64
24 MedRentPctHousInc 2215 non-null float64
25 PctSameCity85     2215 non-null float64
26 PopDens           2215 non-null float64
27 ViolentCrimesPerPop 2215 non-null object
28 nonViolPerPop     2215 non-null object
dtypes: float64(21), int64(6), object(2)
memory usage: 502.0+ KB

```

## Convert Violent and nonViolent Crimes to int

```
In [5]: data[['ViolentCrimesPerPop', 'nonViolPerPop']] = data[['ViolentCrimesPerPop', 'nonViolPerPop']].apply(pd.to_numeric, errors='coerce')
```

```
In [6]: data.isnull().sum()
```

```
Out[6]: householdsize      0
racepctblack            0
racePctWhite            0
racePctAsian            0
racePctHisp             0
agePct12t29             0
pctWWage                0
pctWSocSec              0
pctWPubAsst             0
perCapInc               0
whitePerCap             0
blackPerCap             0
indianPerCap            0
AsianPerCap             0
OtherPerCap             1
HispPerCap              0
PctPopUnderPov          0
PctNotHSGrad            0
PctUnemployed           0
MalePctNevMarr          0
TotalPctDiv             0
PctKidsBornNeverMar     0
PctImmigRec5            0
PctPersDenseHous        0
MedRentPctHousInc       0
PctSameCity85           0
PopDens                 0
ViolentCrimesPerPop     221
nonViolPerPop           97
dtype: int64
```

```
In [7]: # Remove missing values

df = data.dropna()
```

```
In [8]: df.isnull().sum()
```

```
Out[8]: householdsize      0
racepctblack            0
racePctWhite            0
racePctAsian            0
racePctHisp             0
agePct12t29             0
pctWWage                0
pctWSocSec              0
pctWPubAsst             0
perCapInc               0
whitePerCap             0
blackPerCap             0
indianPerCap            0
AsianPerCap             0
OtherPerCap             0
HispPerCap              0
PctPopUnderPov          0
PctNotHSGrad            0
```

```

PctUnemployed      0
MalePctNevMarr     0
TotalPctDiv        0
PctKidsBornNeverMar 0
PctImmigRec5       0
PctPersDenseHous   0
MedRentPctHousInc  0
PctSameCity85      0
PopDens            0
ViolentCrimesPerPop 0
nonViolPerPop      0
dtype: int64


```

```
In [9]: df.describe().T
```

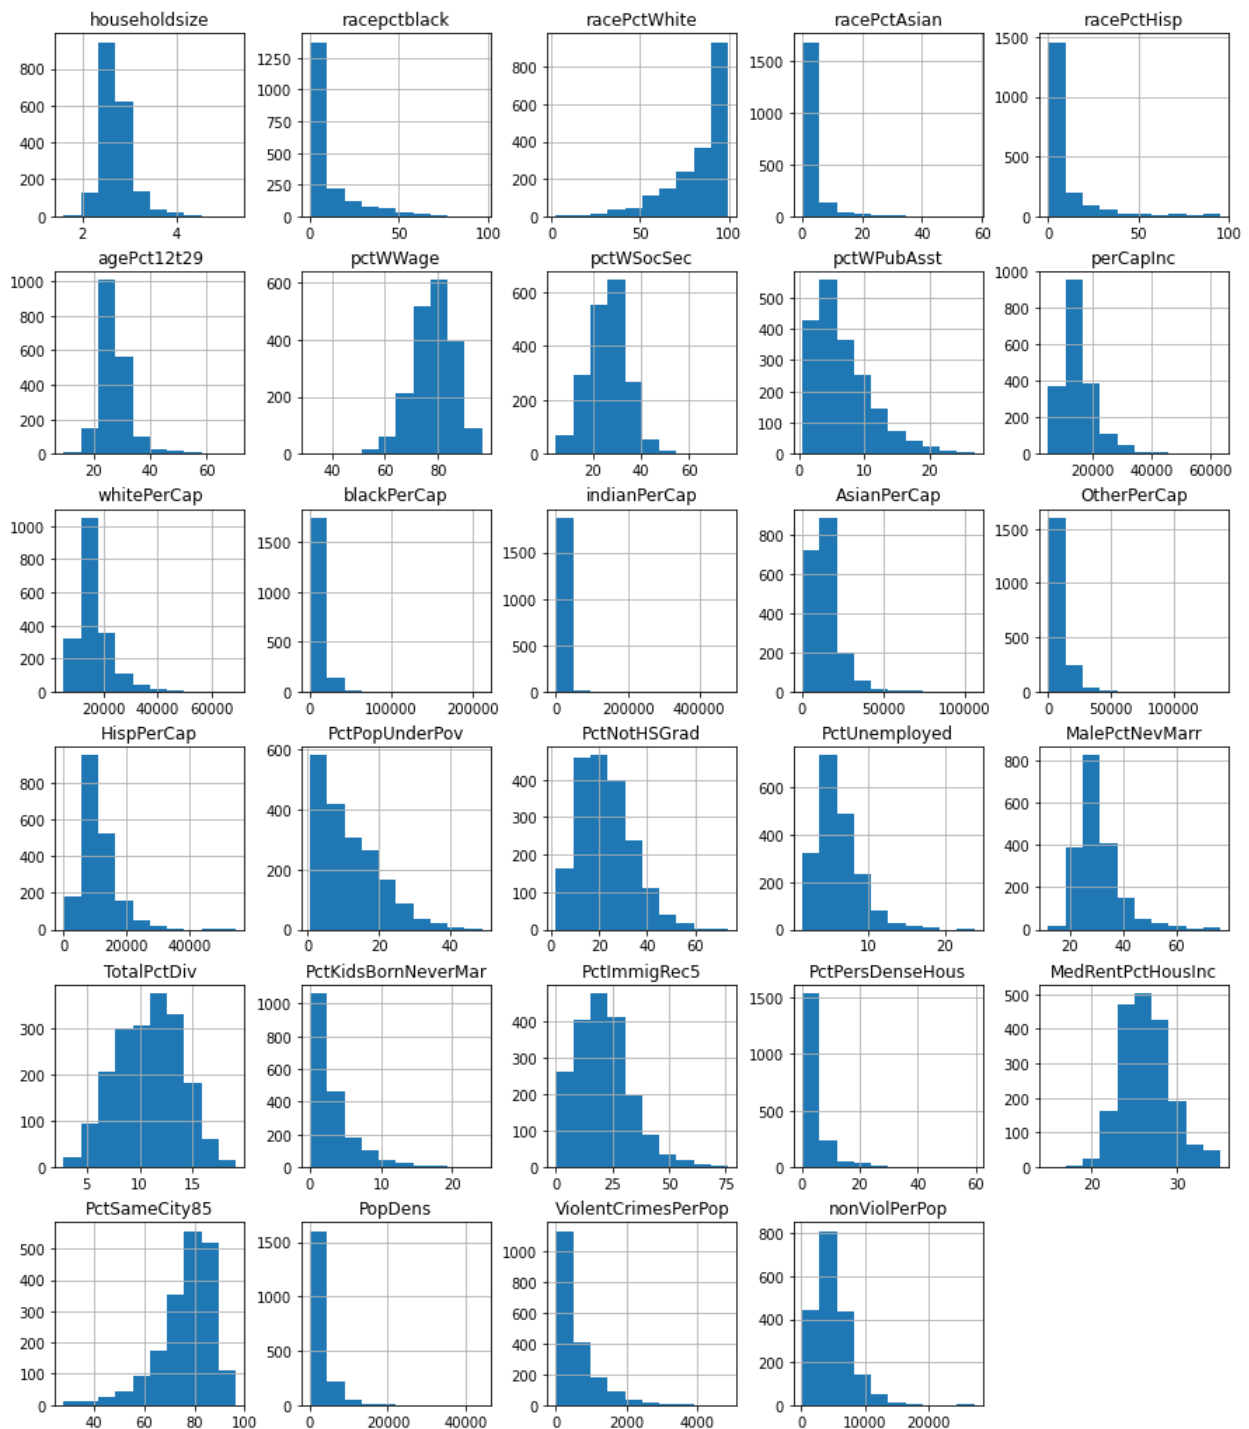
```
Out[9]:
```

|                     | count  | mean         | std          | min     | 25%      | 50%      | 75%      | max       |
|---------------------|--------|--------------|--------------|---------|----------|----------|----------|-----------|
| householdsize       | 1901.0 | 2.712167     | 0.347454     | 1.60    | 2.50     | 2.66     | 2.86     | 5.28      |
| racepctblack        | 1901.0 | 9.358958     | 13.935927    | 0.00    | 0.93     | 3.04     | 11.43    | 96.67     |
| racePctWhite        | 1901.0 | 83.466423    | 16.357057    | 2.68    | 75.77    | 89.61    | 95.96    | 99.63     |
| racePctAsian        | 1901.0 | 2.822799     | 4.738172     | 0.06    | 0.63     | 1.27     | 2.88     | 57.46     |
| racePctHisp         | 1901.0 | 8.718985     | 15.449951    | 0.12    | 0.95     | 2.43     | 8.92     | 95.29     |
| agePct12t29         | 1901.0 | 27.601094    | 6.153583     | 9.38    | 24.37    | 26.78    | 29.20    | 70.51     |
| pctWWage            | 1901.0 | 78.189863    | 7.841396     | 31.68   | 73.45    | 78.55    | 83.76    | 96.62     |
| pctWSocSec          | 1901.0 | 26.577880    | 8.252830     | 4.81    | 20.90    | 26.66    | 31.72    | 76.39     |
| pctWPubAsst         | 1901.0 | 6.755381     | 4.482357     | 0.50    | 3.36     | 5.62     | 9.09     | 26.92     |
| perCapInc           | 1901.0 | 15604.296160 | 6289.037905  | 5237.00 | 11563.00 | 14087.00 | 17910.00 | 63302.00  |
| whitePerCap         | 1901.0 | 16616.188322 | 6381.751863  | 5472.00 | 12643.00 | 15087.00 | 18710.00 | 68850.00  |
| blackPerCap         | 1901.0 | 11582.885324 | 9374.183691  | 0.00    | 6748.00  | 9784.00  | 14549.00 | 212120.00 |
| indianPerCap        | 1901.0 | 12328.749079 | 15519.246128 | 0.00    | 6405.00  | 9943.00  | 14807.00 | 480000.00 |
| AsianPerCap         | 1901.0 | 14293.126775 | 9627.835796  | 0.00    | 8542.00  | 12393.00 | 17351.00 | 106165.00 |
| OtherPerCap         | 1901.0 | 9480.354024  | 8070.968529  | 0.00    | 5615.00  | 8205.00  | 11471.00 | 137000.00 |
| HispPerCap          | 1901.0 | 11036.994740 | 5774.773908  | 0.00    | 7288.00  | 9709.00  | 13431.00 | 54648.00  |
| PctPopUnderPov      | 1901.0 | 11.664277    | 8.467334     | 0.64    | 4.63     | 9.38     | 17.04    | 48.82     |
| PctNotHSGrad        | 1901.0 | 22.656597    | 11.079400    | 2.09    | 14.16    | 21.54    | 29.59    | 73.66     |
| PctUnemployed       | 1901.0 | 6.009932     | 2.705395     | 1.32    | 4.09     | 5.47     | 7.41     | 23.83     |
| MalePctNevMarr      | 1901.0 | 30.654561    | 8.045644     | 12.06   | 25.45    | 29.02    | 33.44    | 76.32     |
| TotalPctDiv         | 1901.0 | 10.867307    | 3.016394     | 2.83    | 8.59     | 11.03    | 13.08    | 19.11     |
| PctKidsBornNeverMar | 1901.0 | 3.109295     | 3.058334     | 0.00    | 1.07     | 2.06     | 3.93     | 24.19     |
| PctImmigRec5        | 1901.0 | 20.784808    | 12.271561    | 0.00    | 11.72    | 19.74    | 27.63    | 76.16     |
| PctPersDenseHous    | 1901.0 | 4.400295     | 5.939846     | 0.15    | 1.31     | 2.49     | 4.96     | 59.49     |
| MedRentPctHousInc   | 1901.0 | 26.352814    | 2.912996     | 14.90   | 24.40    | 26.20    | 28.10    | 35.10     |
| PctSameCity85       | 1901.0 | 76.990142    | 10.837694    | 27.95   | 71.74    | 79.13    | 84.67    | 96.59     |
| PopDens             | 1901.0 | 2804.223461  | 2945.490095  | 10.00   | 1175.60  | 2003.50  | 3278.30  | 44229.90  |
| ViolentCrimesPerPop | 1901.0 | 583.712941   | 608.430184   | 6.64    | 163.75   | 369.30   | 792.93   | 4877.06   |
| nonViolPerPop       | 1901.0 | 4941.049011  | 2786.789842  | 116.79  | 2913.24  | 4479.11  | 6265.54  | 27119.76  |

```
In [10]: corr_matrix = df.corr()
plt.figure(figsize=(15,10))
sns.heatmap(corr_matrix,
            annot=True,
            linewidths=0.5,
            fmt='.2f',
            cmap='YlGnBu');
```



```
In [11]: def show_hist(x):
plt.rcParams["figure.figsize"] = 15,18
x.hist()
show_hist(df)
```



Log Transform the target variables--ViolentCrimesPerPop and nonViolPerPop

```
In [12]: df['Ln_VCrime'] = np.log(df['ViolentCrimesPerPop'])
df['Ln_nVCrime'] = np.log(df['nonViolPerPop'])
```

Drop Violent and nonViolent Crimes variables

```
In [13]: df.drop(['ViolentCrimesPerPop', 'nonViolPerPop'], axis=1, inplace=True)
```

Perform a Baseline algorithm test

```
In [14]: from sklearn.model_selection import train_test_split
```

```
In [15]: X = df.drop(['Ln_VCrime', 'Ln_nVCrime'], axis=1)
y = df['Ln_nVCrime']
```

```
In [16]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
```

## Import the algorithms

```
In [17]: from sklearn.linear_model import LinearRegression, Lasso, ElasticNetCV, ElasticNet
from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
```

We standard the features because some of them are in percentage while some are in real numbers

```
In [18]: pipelines = []
pipelines.append(('ScalerLR', Pipeline([('Scaler', StandardScaler()), ('Lr', LinearRegression())])))
pipelines.append(('ScalerLasso', Pipeline([('Scaler', StandardScaler()), ('LASSO', Lasso())])))
pipelines.append(('ScalerEN', Pipeline([('Scaler', StandardScaler()), ('EN', ElasticNet())])))
pipelines.append(('ScalerKnn', Pipeline([('Scaler', StandardScaler()), ('KNN', KNeighborsRegressor())])))
pipelines.append(('ScalerRf', Pipeline([('Scaler', StandardScaler()), ('RF', RandomForestRegressor())])))

results = []
names = []

for name, model in pipelines:
    kfold = KFold(n_splits=10)
    cv_results = cross_val_score(model, X_train, y_train, cv=kfold, scoring='neg_mean_squared_error')
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
```

```
ScalerLR: -0.154098 (0.038381)
ScalerLasso: -0.335276 (0.048884)
ScalerEN: -0.335276 (0.048884)
ScalerKnn: -0.172971 (0.037560)
ScalerRf: -0.152467 (0.043657)
```

RandomForest Regressor seems to perform better than all the other algorithms, hence we focus on using Randomforest to build the final model

```
In [19]: from sklearn.model_selection import GridSearchCV
```

```
In [20]: from sklearn.feature_selection import SelectKBest, f_regression, VarianceThreshold
```

```
In [21]: pipe = Pipeline([("std", StandardScaler()),
                        ("var", VarianceThreshold()),
                        ("selector", SelectKBest()),
                        ("regressor", RandomForestRegressor())])

params = [
    {"selector_k": [6, 7, 8, 9, 10, 11, 12, 13, 14]},
    {"regressor": [RandomForestRegressor()],
     "regressor__n_estimators": [10, 100, 200, 1000],
     "regressor__max_features": ['auto', 'sqrt', 'log2'],
     "regressor__max_depth": [2, 3, 4, 5, 6, 7]}
]
```

```
In [22]: grid = GridSearchCV(pipe, params, scoring='neg_mean_squared_error', cv=5)
Model_result = grid.fit(X_train, y_train)
```

```
In [23]: best_est = Model_result.best_estimator_
print(best_est)
```

```
Pipeline(steps=[('std', StandardScaler()), ('var', VarianceThreshold()),
                 ('selector', SelectKBest(k=14)),
                 ('regressor', RandomForestRegressor())])
```

```
In [24]: best_est.score(X_test, y_test)
```

```
Out[24]: 0.5798138375912989
```

```
In [26]: y_pred = Model_result.predict(X_test)
y_pred_tr = best_est.predict(X_train)
```

```
r_square = r2_score(y_train,y_pred_tr)
r_square
```

Out[26]: 0.9323748441153882

```
In [27]: r_square_test = r2_score(y_test,y_pred)
r_square_test
```

Out[27]: 0.5798138375912989

```
In [28]: col_after_var = X_train.columns[best_est['var'].get_support()]
mask_sel = best_est['selector'].get_support(indices=True)
final_feature_cols = col_after_var[mask_sel]
```

```
In [29]: final_feature_cols
```

Out[29]: Index(['racePctblack', 'racePctWhite', 'racePctAsian', 'racePctHisp',  
'agePct12t29', 'pctWPubAsst', 'perCapInc', 'whitePerCap',  
'PctPopUnderPov', 'MalePctNevMarr', 'TotalPctDiv',  
'PctKidsBornNeverMar', 'PctPersDenseHous', 'PopDens'],  
dtype='object')

```
In [30]: coef = Model_result.best_estimator_.named_steps['regressor'].feature_importances_
importance = np.abs(coef)
importance
```

Out[30]: array([0.04015851, 0.05860437, 0.0256236 , 0.03619961, 0.04108656,  
0.030938 , 0.0201021 , 0.02916315, 0.24727069, 0.03179994,  
0.27353836, 0.06045644, 0.05103964, 0.05401903])

## Combine the dataframe to know the shortlisted features

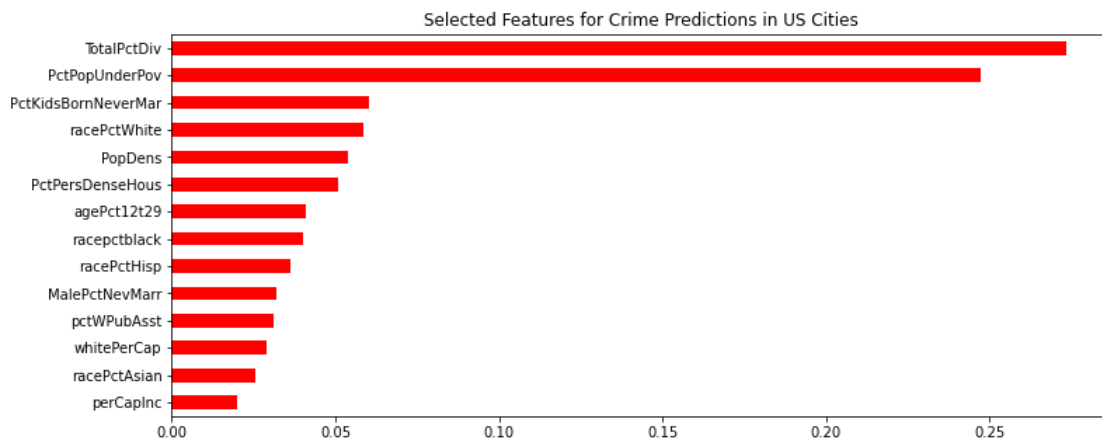
```
In [31]: combination = pd.Series(importance,final_feature_cols)
combination
```

Out[31]:

|                     |          |
|---------------------|----------|
| racePctblack        | 0.040159 |
| racePctWhite        | 0.058604 |
| racePctAsian        | 0.025624 |
| racePctHisp         | 0.036200 |
| agePct12t29         | 0.041087 |
| pctWPubAsst         | 0.030938 |
| perCapInc           | 0.020102 |
| whitePerCap         | 0.029163 |
| PctPopUnderPov      | 0.247271 |
| MalePctNevMarr      | 0.031800 |
| TotalPctDiv         | 0.273538 |
| PctKidsBornNeverMar | 0.060456 |
| PctPersDenseHous    | 0.051040 |
| PopDens             | 0.054019 |

dtype: float64

```
In [32]: combination.sort_values().plot.barh(color='red',figsize=(12,5))
plt.title("Selected Features for Crime Predictions in US Cities");
```



This model identified Total Percentage of People divorced as the number 1 predictor of crimes in the US communities. This was followed by the proportion of the population below poverty line. Again, the proportion of kids born to never married parents influences crime rates. Other factors

include population density,percent of persons in dense housing (more than 1 person per room),proportion of people between age 12 and 29 living in the neighborhood respectively.

Conclusion

There is a need for government to look into how to strengthen families becuae, it seems the major cause of crimes as identified from this exercise is a dysfunctional family.

In [ ]: